

Ансамбли моделей машинного обучения

Цель лабораторной работы: изучение ансамблей моделей машинного обучения. Задание:

1. Выберите набор данных (датасет) для решения задачи классификации или регрессии.
2. В случае необходимости проведите удаление или заполнение пропусков и кодирование категориальных признаков.
3. С использованием метода `train_test_split` разделите выборку на обучающую и тестовую.
4. Обучите две ансамблевые модели. Оцените качество моделей с помощью одной из подходящих для задачи метрик. Сравните качество полученных моделей.
5. Произведите для каждой модели подбор значений одного гиперпараметра. В зависимости от используемой библиотеки можно применять функцию `GridSearchCV`, использовать перебор параметров в цикле, или использовать другие методы.
6. Повторите пункт 4 для найденных оптимальных значений гиперпараметров. Сравните качество полученных моделей с качеством моделей, полученных в пункте 4.

In [1]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
sns.set(style="ticks")
data = pd.read_csv('Data/lab_4/heart.csv', sep=",")
data.head(5)
```

Out[1]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	t
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	

In [2]:

```
data.shape
```

Out[2]:

(303, 14)

In [3]:

```
# Проверка на пустые значения
data.isnull().sum()
```

Out[3]:

```
age          0
sex          0
cp           0
trestbps     0
chol         0
fbs          0
restecg      0
thalach      0
exang        0
oldpeak      0
slope        0
ca           0
thal         0
target       0
dtype: int64
```

Feature Scaling

In [4]:

```
from sklearn.preprocessing import MinMaxScaler
import warnings
warnings.filterwarnings('ignore')

# Create the scaler object with a range of 0-1
scaler = MinMaxScaler(feature_range=(0, 1))
# Fit on data, transform data
scaler.fit_transform(data)
```

Out[4]:

```
array([[0.70833333, 1.          , 1.          , ..., 0.          , 0.33333333,
        1.          ],
       [0.16666667, 1.          , 0.66666667, ..., 0.          , 0.66666667,
        1.          ],
       [0.25       , 0.          , 0.33333333, ..., 0.          , 0.66666667,
        1.          ],
       ...,
       [0.8125     , 1.          , 0.          , ..., 0.5       , 1.          ,
        0.          ],
       [0.58333333, 1.          , 0.          , ..., 0.25      , 1.          ,
        0.          ],
       [0.58333333, 0.          , 0.33333333, ..., 0.25      , 0.66666667,
        0.          ]])
```

In [5]:

```
from sklearn import svm
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import roc_curve, auc
import pylab as pl
```

In [6]:

```
# Пустых значений нет
# Перейдем к разделению выборки на обучающую и тестовую.
X = data.drop('target', axis = 1).values
y = data['target'].values
```

Ансамблевые модели

In [7]:

```
from sklearn.model_selection import train_test_split
kfold = 5 #количество подвыборок для валидации
itog_val = {} #список для записи результатов кросс валидации разных алгоритмов
```

In [8]:

```
ROctrainTRN, ROctestTRN, ROctrainTRG, ROctestTRG = train_test_split(X, y, test_size=0.2
5)
```

In [9]:

```
model_rfc = RandomForestClassifier(n_estimators = 70) #в параметре передаем кол-во дере
вьев
model_knc = KNeighborsClassifier(n_neighbors = 18) #в параметре передаем кол-во соседей
model_lr = LogisticRegression(penalty='l1', tol=0.01)
model_svc = svm.SVC() #по умолчанию kernel='rbf'
```

In [10]:

```
from sklearn.model_selection import cross_val_score
```

SVM - метод опорных векторов(SVC): Суть работы “Машин” Опорных Векторов проста: алгоритм создает линию или гиперплоскость, которая разделяет данные на классы. Метод k-ближайших соседей(KNeighborsClassifier) Random forest(RandomForestClassifier) Логистическая регрессия (LogisticRegression)

In [11]:

```
from sklearn.ensemble import ExtraTreesClassifier, ExtraTreesRegressor
```

In [12]:


```
scores = cross_val_score(model_rfc, X, y, cv = kfold)
itog_val['RandomForestClassifier'] = scores.mean()
scores = cross_val_score(model_knc, X, y, cv = kfold)
itog_val['KNeighborsClassifier'] = scores.mean()
scores = cross_val_score(model_lr, X, y, cv = kfold)
itog_val['LogisticRegression'] = scores.mean()
scores = cross_val_score(model_svc, X, y, cv = kfold)
itog_val['SVC'] = scores.mean()
```

In [13]:

```
import matplotlib.pyplot as plt
plt.style.use('bmh')
data.from_dict(data = itog_val, orient='index').plot(kind='bar', legend=False)
```

Out[13]:

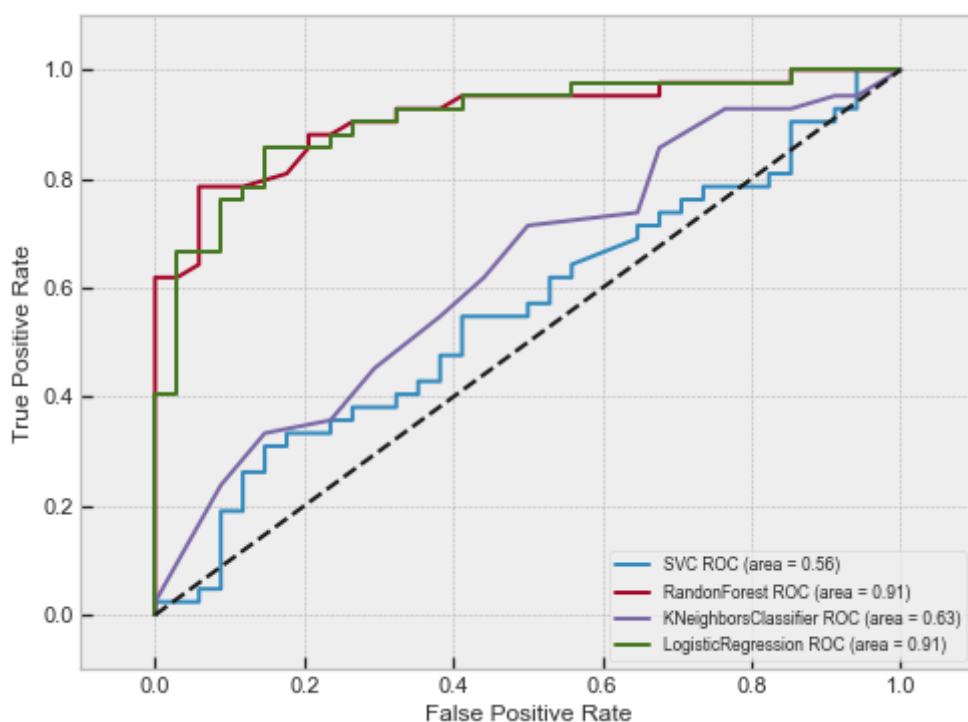
<matplotlib.axes._subplots.AxesSubplot at 0x1a1e343240>



In [14]:

```
pl.clf()
plt.figure(figsize=(8,6))
#SVC
model_svc.probability = True
probas = model_svc.fit(ROctrainTRN, ROctrainTRG).predict_proba(ROctestTRN)
fpr, tpr, thresholds = roc_curve(ROctestTRG, probas[:, 1])
roc_auc = auc(fpr, tpr)
pl.plot(fpr, tpr, label='%s ROC (area = %0.2f)' % ('SVC', roc_auc))
#RandomForestClassifier
probas = model_rfc.fit(ROctrainTRN, ROctrainTRG).predict_proba(ROctestTRN)
fpr, tpr, thresholds = roc_curve(ROctestTRG, probas[:, 1])
roc_auc = auc(fpr, tpr)
pl.plot(fpr, tpr, label='%s ROC (area = %0.2f)' % ('RandomForest',roc_auc))
#KNeighborsClassifier
probas = model_knc.fit(ROctrainTRN, ROctrainTRG).predict_proba(ROctestTRN)
fpr, tpr, thresholds = roc_curve(ROctestTRG, probas[:, 1])
roc_auc = auc(fpr, tpr)
pl.plot(fpr, tpr, label='%s ROC (area = %0.2f)' % ('KNeighborsClassifier',roc_auc))
#LogisticRegression
probas = model_lr.fit(ROctrainTRN, ROctrainTRG).predict_proba(ROctestTRN)
fpr, tpr, thresholds = roc_curve(ROctestTRG, probas[:, 1])
roc_auc = auc(fpr, tpr)
pl.plot(fpr, tpr, label='%s ROC (area = %0.2f)' % ('LogisticRegression',roc_auc))
pl.plot([0, 1], [0, 1], 'k--')
pl.xlim([-0.1, 1.1])
pl.ylim([-0.1, 1.1])
pl.xlabel('False Positive Rate')
pl.ylabel('True Positive Rate')
pl.legend(loc=0, fontsize='small')
pl.show()
```

<Figure size 432x288 with 0 Axes>



In [15]:

```
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier
from sklearn.metrics import accuracy_score
from sklearn.metrics import balanced_accuracy_score
from sklearn.metrics import precision_score, recall_score, f1_score
```

In [16]:

```
from sklearn.model_selection import train_test_split
# Функция train_test_split разделила исходную выборку таким образом,
# чтобы в обучающей и тестовой частях сохранились пропорции классов.
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.35, random_state=1)
```

In [17]:

```
from sklearn.preprocessing import MinMaxScaler
import warnings
warnings.filterwarnings('ignore')

# Create the scaler object with a range of 0-1
scaler = MinMaxScaler(feature_range=(0, 1))
# Fit on data, transform data
scaler.fit_transform(X)
scaler.fit_transform(X_train)
scaler.fit_transform(X_test)
```

Out[17]:

```
array([[0.77777778, 0.          , 0.          , ..., 0.          , 0.75        ,
        1.          ],
       [0.61111111, 1.          , 0.33333333, ..., 1.          , 0.          ,
        1.          ],
       [0.38888889, 1.          , 0.          , ..., 1.          , 0.5         ,
        1.          ],
       ...,
       [0.52777778, 1.          , 0.          , ..., 0.          , 0.          ,
        1.          ],
       [0.66666667, 1.          , 0.33333333, ..., 0.5         , 1.          ,
        1.          ],
       [0.38888889, 1.          , 0.33333333, ..., 0.          , 0.          ,
        1.          ]])
```

In [18]:

```
# n_estimators = 10 (default)
rfc = RandomForestClassifier().fit(X_train, y_train)
predicted_rfc = rfc.predict(X_test)
```

In [19]:

```
accuracy_score(y_test, predicted_rfc)
```

Out[19]:

```
0.7570093457943925
```

In [20]:

```
balanced_accuracy_score(y_test, predicted_rfc)
```

Out[20]:

0.7547368421052632

In [21]:

```
(precision_score(y_test, predicted_rfc, average='weighted'),  
 recall_score(y_test, predicted_rfc, average='weighted'))
```

Out[21]:

(0.7567717408522097, 0.7570093457943925)

In [22]:

```
f1_score(y_test, predicted_rfc, average='weighted')
```

Out[22]:

0.7566245963419206

In [23]:

```
# n_estimators = 50 (default)  
abc = AdaBoostClassifier().fit(X_train, y_train)  
predicted_abc = abc.predict(X_test)
```

In [24]:

```
accuracy_score(y_test, predicted_abc)
```

Out[24]:

0.7289719626168224

In [25]:

```
balanced_accuracy_score(y_test, predicted_abc)
```

Out[25]:

0.7284210526315789

In [26]:

```
(precision_score(y_test, predicted_abc, average='weighted'),  
 recall_score(y_test, predicted_abc, average='weighted'))
```

Out[26]:

(0.7293842770753162, 0.7289719626168224)

In [27]:

```
f1_score(y_test, predicted_abc, average='weighted')
```

Out[27]:

0.7291144464706996

In [28]:

```
rfc_n_range = np.array(range(5,100,5))
rfc_tuned_parameters = [{'n_estimators': rfc_n_range}]
rfc_tuned_parameters
```

Out[28]:

```
[{'n_estimators': array([ 5, 10, 15, 20, 25, 30, 35, 40, 45, 50, 55, 60, 65, 70, 75, 80, 85, 90, 95])}]
```

In [29]:

```
import warnings
from sklearn.model_selection import GridSearchCV
warnings.filterwarnings('ignore')

gs_rfc = GridSearchCV(RandomForestClassifier(), rfc_tuned_parameters, cv=5,
                      scoring='accuracy')
gs_rfc.fit(X_train, y_train)
```

Out[29]:

```
GridSearchCV(cv=5, error_score='raise-deprecating',
             estimator=RandomForestClassifier(bootstrap=True, class_weight=None,
             criterion='gini',
             max_depth=None, max_features='auto', max_leaf_nodes=None,
             min_impurity_decrease=0.0, min_impurity_split=None,
             min_samples_leaf=1, min_samples_split=2,
             min_weight_fraction_leaf=0.0, n_estimators='warn', n_jobs=None,
             oob_score=False, random_state=None, verbose=0,
             warm_start=False),
             fit_params=None, iid='warn', n_jobs=None,
             param_grid=[{'n_estimators': array([ 5, 10, 15, 20, 25, 30, 35, 40,
             45, 50, 55, 60, 65, 70, 75, 80, 85,
             90, 95])}],
             pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
             scoring='accuracy', verbose=0)
```

In [30]:

```
gs_rfc.best_params_
```

Out[30]:


```
{'n_estimators': 85}
```


In [31]:

```
plt.plot(rfc_n_range, gs_rfc.cv_results_['mean_test_score'])
```

Out[31]:

[<matplotlib.lines.Line2D at 0x1a1e9302e8>]



In [32]:

```
abc_n_range = np.array(range(5,100,5))
abc_tuned_parameters = [{'n_estimators': abc_n_range}]
abc_tuned_parameters
```

Out[32]:

```
[{'n_estimators': array([ 5, 10, 15, 20, 25, 30, 35, 40, 45, 50, 55, 60, 65, 70, 75, 80, 85, 90, 95])}]
```

In [33]:

```
gs_abc = GridSearchCV(AdaBoostClassifier(), abc_tuned_parameters, cv=5,
                      scoring='accuracy')
gs_abc.fit(X_train, y_train)
```

Out[33]:

```
GridSearchCV(cv=5, error_score='raise-deprecating',
             estimator=AdaBoostClassifier(algorithm='SAMME.R', base_estimator=None,
             learning_rate=1.0, n_estimators=50, random_state=None),
             fit_params=None, iid='warn', n_jobs=None,
             param_grid=[{'n_estimators': array([ 5, 10, 15, 20, 25, 30, 35, 40, 45, 50, 55, 60, 65, 70, 75, 80, 85, 90, 95])}],
             pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
             scoring='accuracy', verbose=0)
```

In [34]:

```
gs_abc.best_params_
```

Out[34]:


```
{'n_estimators': 25}
```

In [35]:

```
plt.plot(abc_n_range, gs_abc.cv_results_['mean_test_score'])
```

Out[35]:

[<matplotlib.lines.Line2D at 0x1a1e985d68>]



In [36]:

```
rfc_optimized = RandomForestClassifier(n_estimators=gs_rfc.best_params_['n_estimators'])  
.fit(X_train, y_train)  
predicted_rfc_opt = rfc_optimized.predict(X_test)
```

In [37]:

```
accuracy_score(y_test, predicted_rfc_opt)
```

Out[37]:

0.7663551401869159

In [38]:

```
balanced_accuracy_score(y_test, predicted_rfc_opt)
```

Out[38]:

0.7635087719298246

In [39]:

```
(precision_score(y_test, predicted_rfc_opt, average='weighted'),  
 recall_score(y_test, predicted_rfc_opt, average='weighted'))
```

Out[39]:

(0.7663352555179958, 0.7663551401869159)

In [40]:

```
f1_score(y_test, predicted_rfc_opt, average='weighted')
```

Out[40]:

0.765737522265126

In [41]:

```
abc_optimized = RandomForestClassifier(n_estimators=gs_abc.best_params_['n_estimators'])  
.fit(X_train, y_train)  
predicted_abc_opt = abc_optimized.predict(X_test)
```

In [42]:

```
accuracy_score(y_test, predicted_abc_opt)
```

Out[42]:

0.7476635514018691

In [43]:

```
balanced_accuracy_score(y_test, predicted_abc_opt)
```

Out[43]:

0.7471929824561403

In [44]:

```
(precision_score(y_test, predicted_abc_opt, average='weighted'),  
 recall_score(y_test, predicted_abc_opt, average='weighted'))
```

Out[44]:

(0.748059504175502, 0.7476635514018691)

In [45]:

```
f1_score(y_test, predicted_abc_opt, average='weighted')
```

Out[45]:

0.7477962087830651